

Generative Adversarial Networks

Team Toutpourelagan

Lucas Henneçon
Ilian Benaissa-Lejay
Simon Liétar

Wasserstein GAN

- Use Wasserstein distance instead of Jensen-Shannon divergence:

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

- Using the Kantorovich-Rubinstein duality [1]:

$$W(\mathbb{P}_r, \mathbb{P}_\theta) = \sup_{\|f\|_L \leq 1} \mathbb{E}_{x \sim \mathbb{P}_r} [f(x)] - \mathbb{E}_{x \sim \mathbb{P}_\theta} [f(x)]$$

- Loss function of the WGAN:

$$\min_G \max_{\|D\|_L \leq 1} \mathbb{E}_{x \sim \mathbb{P}_r} [D(x)] - \mathbb{E}_{x \sim \mathbb{P}_G} [D(x)]$$

⇒ Discriminator 1 Lipschitz continuous

Wasserstein GAN

- Don't use a sigmoid at the output of D : output not in $[0,1]$
- To ensure Lipschitz-continuity of D, clip the weights between $[-c, c]$

Algorithm:

Update of D (nd times):

- Sample (X_i) batch of real images,
- Generate batch (Z_i) with G
- Backprop and RMSProp as optimizer
- Clip weights between $[-c, c]$

$$\text{Loss} = \frac{1}{N} \sum_{i=1}^N D(Z_i) - \frac{1}{N} \sum_{i=1}^N D(X_i)$$

Update of G:


- Generate batch (Z_i) with G
- Backprop and RMSProp as optimizer

$$\text{Loss} = -\frac{1}{N} \sum_{i=1}^N [D(Z_i)]$$


Wasserstein GAN

- Hyperparameters from the paper
- FID = 70

Generated Images



Discriminator Loss over Epochs



Hyperparameters = 5e-5, batch_size=64, critic_iterations=5, weight_clip =0.01

Wasserstein GAN - Gradient Penalty

- New way to ensure 1-Lipschitz continuity of D [2]:

$$L = \underbrace{\mathbb{E}_{\tilde{x} \sim \mathbb{P}_G}[D(\tilde{x})] - \mathbb{E}_{x \sim \mathbb{P}_r}[D(x)]}_{\text{Original WGAN Discriminator loss}} + \lambda \underbrace{\mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}} [(\|\nabla_{\hat{x}} D(\hat{x})\|_2 - 1)^2]}_{\text{Gradient penalty}}.$$

Algorithm:

Update of D (nd times):



- Sample (X_i) batch of real images, Generate batch (Z_i) with G
- $Q_i = \epsilon X_i + (1-\epsilon)Z_i$ with ϵ following $U[0,1]$
- $\text{Loss} = \frac{1}{N} \sum_{i=1}^N D(Z_i) - \frac{1}{N} \sum_{i=1}^N D(X_i) + \lambda \sum_{i=1}^N (\|\nabla_{Q_i} D(Q_i)\|_2 - 1)^2$
- Backprop and Adam as optimizer

Update of G:

- Same as WGAN with Adam as optimizer

Wasserstein GAN-GP

- More stable
- FID = 35.5, Precision = 0.45, Recall = 0.27



What we'll try next

- Rejection Sampling (find a criterion)
- Fine tuning
- fGANS

Computing precision and recall

- Based on nearest-neighbor search as introduced by Kynkäänniemi et al. (2019) and implemented using a k-d tree
- The paper computes precision and recall using an intermediate layer of VGG-16, which may not be suitable here

